

The Computational Linguistics Summarization Pilot Task @ TAC 2014

Kokil Jaidka*, Muthu Kumar Chandrasekaran, Rahul Jha,
Christopher Jones, Min-Yen Kan, Ankur Khanna, Diego Molla-
Aliod, Dragomir R. Radev, Francesco Ronzano, Horacio Saggion

* Nanyang Technological University, Singapore
koki0001@e.ntu.edu.sg

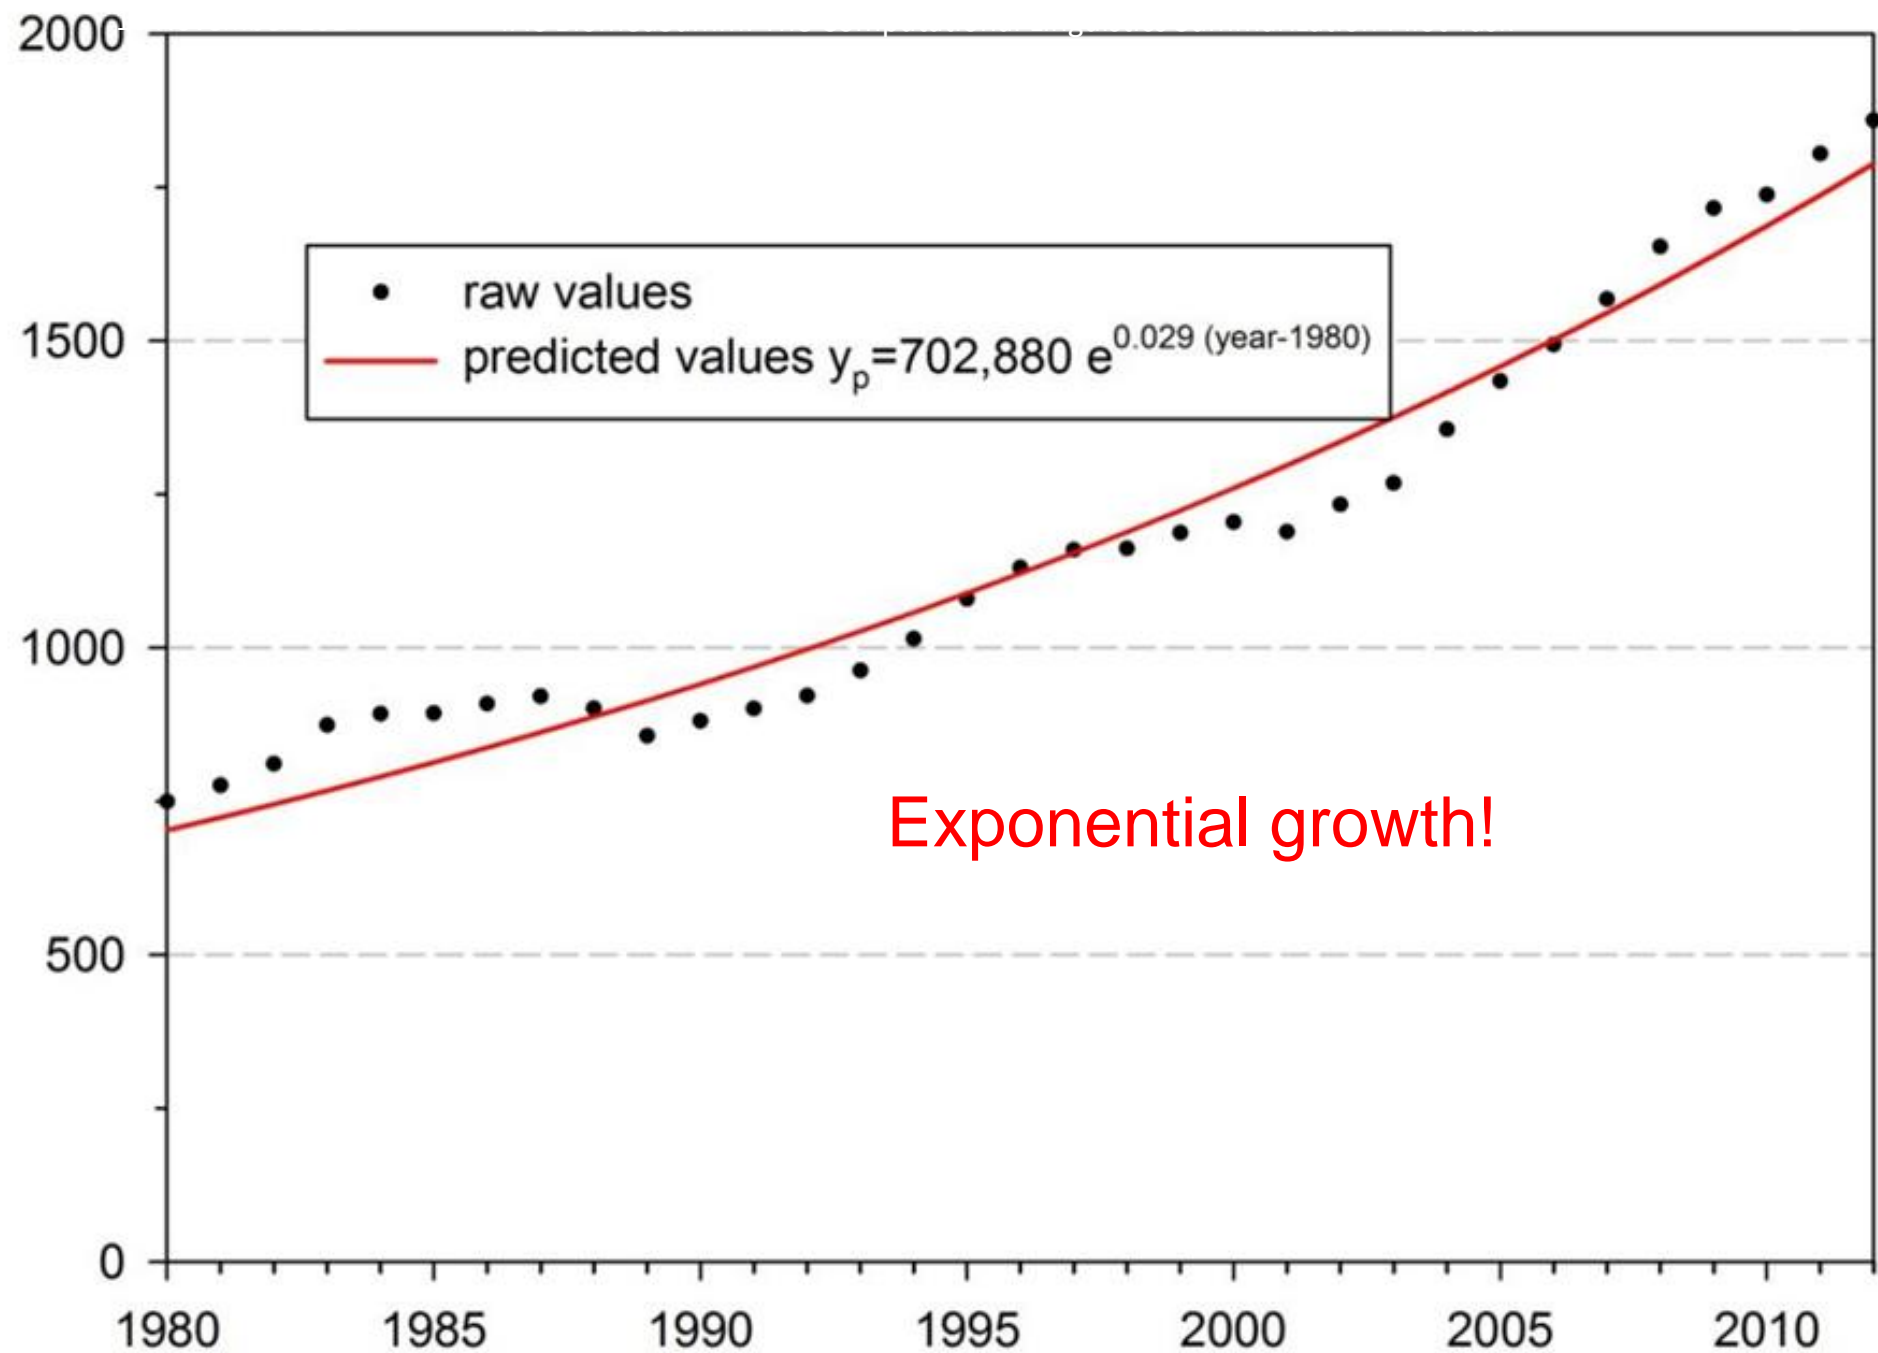


Scientific Document Summarization

I have an abstract

I am done!

[Photo Credits Dennis Jarvis @flickr](#)



Outline

- CL-Summ so far
 - Citation based extractive summaries
 - Faceted summaries
 - Automatic literature review
 - ACL corpus
- The CL-Summ Shared Task
- TAC 2015: CL-Summ track
- Acknowledgements

Scientific Document Summarization

- Abstract
 - Authors' own summary
- Citation summary
 - Community creates a summary when citing
- Faceted summary
 - Capture all aspects of a paper

Towards Automated Related Work Summarization

Cong Duy Vu Hoang and **Min-Yen Kan**

Department of Computer Science

School of Computing

National University of Singapore

{hcdvu, kanmy}@comp.nus.edu.sg

Abstract

We introduce the novel problem of automatic related work summarization. Given multiple articles (*e.g.*, conference/journal papers) as input, a related work summarization system creates a topic-biased summary of related work specific to the target paper. Our prototype **Related Work Summarization** system, **ReWoS**, takes in set of keywords arranged in a hierarchical fashion that describes a target paper's topics, to drive the creation of an extractive summary using two different strategies for locating appropriate sentences for general topics as well as detailed ones. Our initial

Framework for Scientific document summarization

We dissect the full challenge as bringing together work of disparate interests; 1) in finding relevant documents; 2) in identifying the salient aspects of these documents in relation to the current work worth summarizing; and 3) in generating the final topic-biased summary. While it is clear that current NLP technology does not let us build a complete solution for this task, we believe that tackling the individual components will help bring us towards an eventual solution.

Using Citations to Generate Surveys of Scientific Paradigms

**Saif Mohammad^{†*}, Bonnie Dorr^{†‡*}, Melissa Egan^{†‡}, Ahmed Hassan^φ,
Pradeep Muthukrishnan^φ, Vahed Qazvinian^φ, Dragomir Radev^{φ§}, David Zajic^{‡◇}**

Institute for Advanced Computer Studies[†] and Computer Science[‡], University of Maryland.
Human Language Technology Center of Excellence^{*}. Center for Advanced Study of Language[◇].
{saif,bonnie,mkegan,dmzajic}@umiacs.umd.edu

Department of Electrical Engineering and Computer Science^φ
School of Information[§], University of Michigan.
{hassanam,mpradeep,vahed,radev}@umich.edu

Citations 1) select papers and
2) identify salient parts of the
cited paper.

first steps in producing an automatically generated, readily consumable, technical survey. Specifically we explore the combination of citation information and summarization techniques. Even though prior work (Teufel et

some of which may be unfamiliar to panelists. Thus, they must learn about a new discipline “on the fly” in order to relate their own expertise to the proposal.

Our goal is to effectively serve these needs by combining two currently available technologies: (1) bibliometric lexical link mining that exploits the structure of citations and relations among citations; and (2) summarization techniques that exploit the content of the material in both the citing and cited papers.

It is generally agreed upon that manually written

Scientific Document Summarization

Citation based extractive summaries

Scope of Citation


- Qazvinian, V., and Radev, D. R. “Identifying non-explicit citing sentences for citation-based summarization” (ACL, 2010)
- Abu-Jbara, Amjad, and Dragomir Radev. “Reference scope identification in citing sentences.” (ACL, 2012)

Citation summary

Blind Men and Elephants: What Do Citation Summaries Tell Us About a Research Article?

Aaron Elkiss, Siwei Shen, Anthony Fader, Güneş Erkan, David States, and Dragomir Radev
University of Michigan, Ann Arbor MI 48109

The old Asian legend about the blind men and the elephant comes to mind when looking at how different authors of scientific papers describe a piece of related prior work. It turns out that different citations to the same paper often focus on different aspects of that paper and that neither provides a full description of its full set of contributions. In this article, we will describe our investigation of this phenomenon. We studied cita-



Citations are like celebrity quotes taken out of context

[Image credits Ken Ammi @flickr](#)

Deconstructing Human Literature Reviews – A Framework for Multi-Document Summarization

Kokil Jaidka, Christopher S.G. Khoo, Jin-Cheon Na
Division of Information Studies

Wee Kim Wee School of Communication and Information
Nanyang Technological University, Singapore

[kokil, chriskhoo]@pmail.ntu.edu.sg, tjcna@ntu.edu.sg

Scientific summaries are
written to fulfill information
& argumentative functions

- Identify new research questions
- Define the proposed research contributions
- Build the justification for the current work
- Situate the work in the research literature
- Reinterpret and critique previous results

These rhetorical characteristics of literature reviews make it a challenging research problem in automatic multi-document summarization – not only should the summarizer identify salient information, but it should also synthesize the summary in a way that achieves certain argumentative purposes. The problem of summarization in context was first identified by Sparck Jones and Endres-Niggemeyer (1995) and subsequently in Sparck Jones' follow up article

Scientific Document Summarization

Citation-based extractive summaries

Scope of Citation

- Qazvinian, V., and Radev, D. R. “Identifying non-explicit citing sentences for citation-based summarization” (ACL, 2010)
- Abu-Jbara, Amjad, and Dragomir Radev. “Reference scope identification in citing sentences.” (ACL, 2012)

Coherence

- Abu-Jbara, Amjad, and Dragomir Radev. “Coherent citation-based summarization of scientific papers.” (ACL 2011)

ABSTRACT

Objective: Obesity and disturbed eating behaviors are both associated with low self-esteem and poor body images. The evolution of men's body images before and after a weight loss program were evaluated.

Methods: In total, 100 obese men participated in a weight loss program lasting 18 weeks, addressing health and body image.

Results: Obese men had a lower quality of life ($p < 0.001$) and a poorer body image ($p < 0.001$).

Conclusion: DTW allowed obese patients to reset both their somatic and psychic constraints on their body image.

Practice implications: Patients are usually reluctant to practice physical activity. Dance therapy addresses not only body image, but also psycho-social aspects of their personality.

Faceted summaries are structured abstracts.

Common in domains such as
Medicine, Biomedical, Bioinformatics

Summarising Scientific Articles — Experiments with Relevance and Rhetorical Status

Simone Teufel*
Cambridge University

Marc Moens*
Rhetorical Systems
and University of Edinburgh

In this paper we propose a strategy for the summarisation of scientific articles which concentrates on the rhetorical status of statements in the article: material for summaries is selected in such a way that summaries can highlight the new contribution of the source paper and situate it with respect to earlier work.

Argumentative zones as facets demarcate new contributions of a paper from background work.

In summary,

- Community concurs that a citation based summary of a scientific document is worthwhile to create
- Assigning facets help create coherent summaries
- Facets that correspond to structure of a scientific documents help us do this



All Fields ▾

Search...

Search Q

Login

Bookmarks | History

Oct 2014: The Third CIPS-SIGHAN Joint Conference on Chinese Language Processing has now been added to the Anthology. The Proceedings of Traitement of Automatique des Langues Naturelles 2014 and its five associated workshops are now available on the Anthology.

Welcome to the ACL Anthology

The ACL Anthology currently hosts 33846 papers on the study of computational linguistics and natural language processing. [Subscribe to the mailing list](#) to receive announcements and updates to the Anthology.

ACL Events	Present - 2010	2009 - 2000	1999 - 1990	1989 - 1974
CL	14 13 12 11 10	09 08 07 06 05 04 03 02 01 00	99 98 97 96 95 94 93 92 91 90	89 88 87 86 85 84 83 82 81 80 79-74
TACL	14 13			
ACL	14 13 12 11 10	09 08 07 06 05 04 03 02 01 00	99 98 97 96 95 94 93 92 91 90	89 88 87 86 85 84 83 82 81 80 79
EACL	14 12	09 06 03	99 97 95 93 91	89 87 85 83
NAACL	13 12 10	09 07 06 04 03 01 00		
*SEMEVAL	14 13 12 10	07 04 01	98	
ANLP		00	97 94 92	88 83
EMNLP	14 13 12 11 10	09 08 07 06 05 04 03 02 01 00	99 98 97 96	
CONLL	14 13 12 11 10	09 08 07 06 05 04 03 02 01 00	99 98 97	
WS	14 13 12 11 10	09 08 07 06 05 04 03 02 01 00	99 98 97 96 95 94 93 91 90	
SIGs	ANN BIOMED DAT DIAL FSM GEN HAN HUM LEX MEDIA MOL MT NLL PARSE MORPHON SEM SEMITIC SLPAT WAC			

Non-ACL Events	Present - 2010	2009 - 2000	1999 - 1990	1989 - 1974	1979 - 1965
COLING	14 12 10	08 06 04 02 00	98 96 94 92 90	88 86 82 80	73 69 67 65
HLT	13 12 10	09 08 07 06 05 04 03 01	94 93 92 91 90	89 86	
IJCNLP	13 11	09 08 05			
LREC	14 12 10	08 06 04 02 00			
PACLIC	13 12 11 10	09 08 07 06 05 04 03 02 01 00	99 98 96 95		
ROCLING	13 12 11 10	09 08 07 06 05 04 03 02 01 00	99 98 97 96 95 94 93 92 91 90	89 88	
TINLAP				87	78 75
ALTA	13 12 11 10	09 08 07 06 05 04 03			
RANLP	13 11	09			

Outline

- CL-Summ ground work done so far
- The CL-Summ corpus, task evaluation
 - Highlights
 - Annotation
 - Evaluation results
- TAC 2015: CL-Summ track
- Acknowledgements

CL-Summ Pilot: highlights

First corpus in the computational linguistics community incorporating prior research on citation based summaries

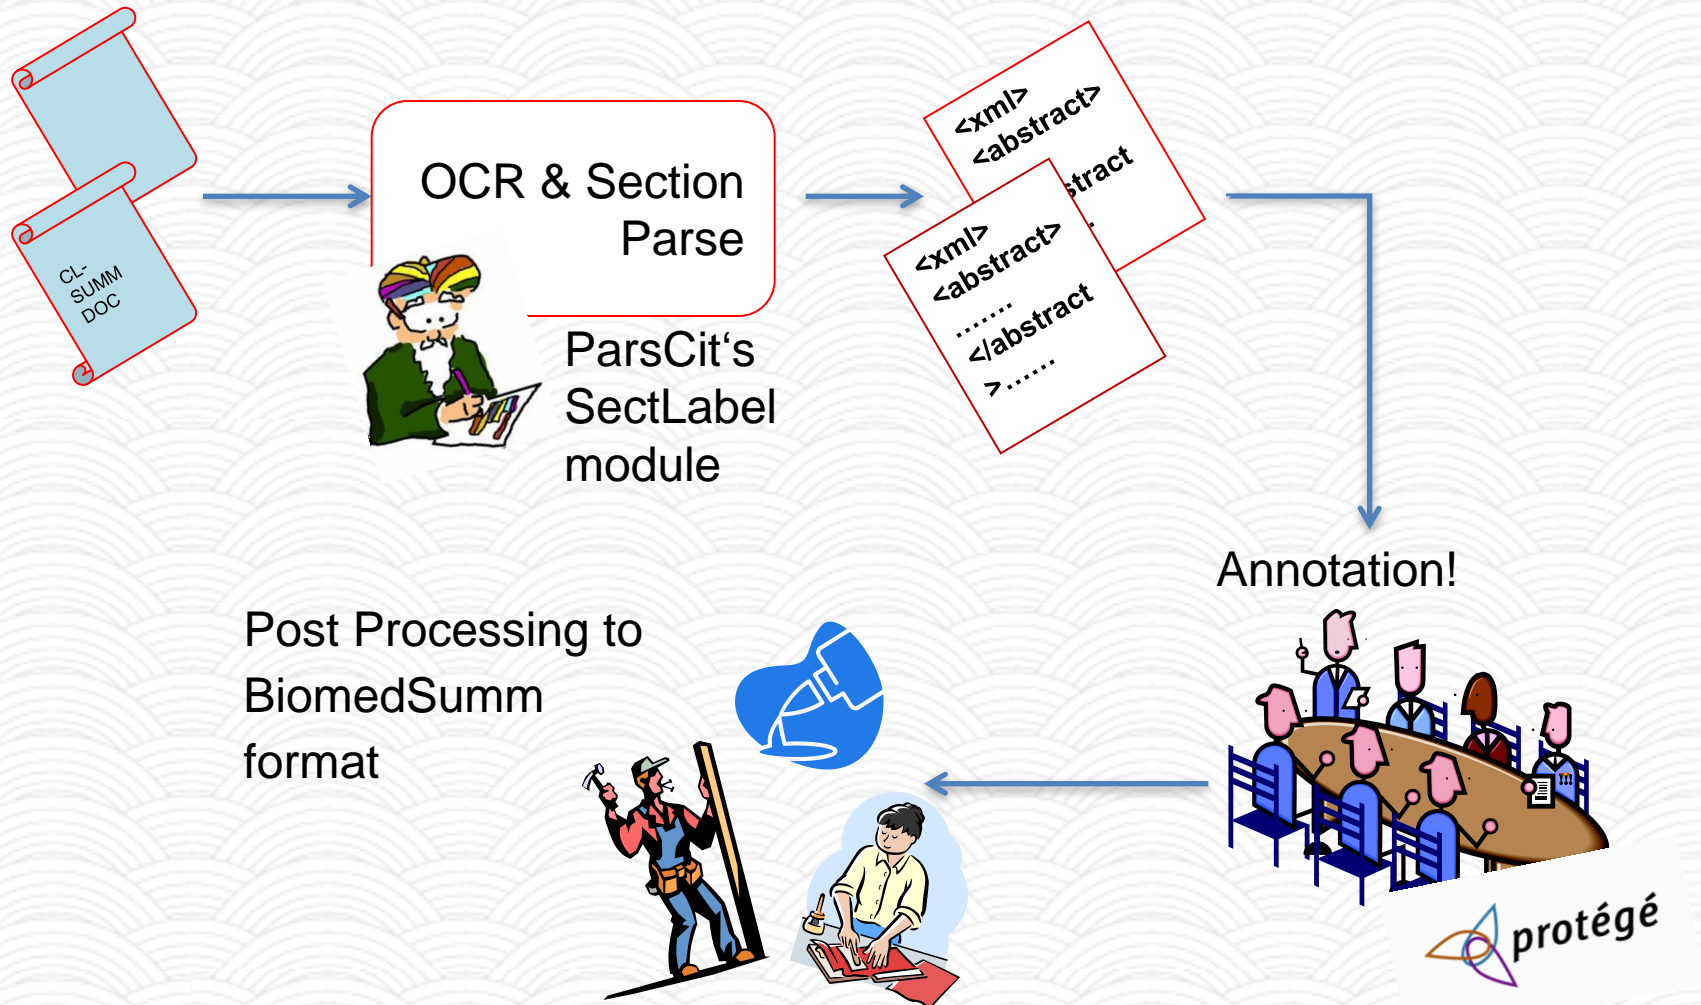
- 10 teams registered
- 3 teams participated in the evaluation
- 2 teams submitted their systems' performance
- 1 more proposed algorithms to solve the tasks

CL training Corpus

- 10 reference papers or topics randomly sampled from the ACL live anthology
- Up to 10 citing papers per reference paper including those outside ACL live anthology
- Annotated corpus publicly available


<https://github.com/WING-NUS/scisumm-corpus/>

Annotation Pipeline




Annotating the SciSumm corpus

- 3 annotators
- Released data has one gold standard annotation per topic or reference paper
- Discourse facet has a minor change from Biomedsumm's categories



Tasks

Task 1A: Identify the text span in the RP which corresponds to the *citances* from the CP.




Tasks

Task 1B: Identify the discourse facet for every cited text span from a predefined set of facets.


Classify the
cited text in
RP into one
of several
facets

Task 1B



Tasks

Task 2: Generate a faceted summary of up to 250 words, of the reference paper, using itself and the citing papers.



Use citances
and the RP to
create a
summary

Evaluation

Small corpus: 10 fold cross validated evaluation over the 10 documents

- Task 1A scored by ROUGE-L metric
- Task 1B scored by classification metrics: Precision, Recall and F_1
- Task 2 also scored by ROUGE-L metric

Results – Task 1A

MQ			Clair_UMich		
Precision	Recall	F_1	Precision	Recall	F_1
0.212	0.335	0.223	0.444	0.574	0.487

- MQ was unsupervised while Clair_Umich was supervised
- Challenging classification problem: Task seeks to map each citation sentence with a few out of 100s of potential matches in the Reference paper (RP)
- Lexical, semantic and structural similarities between citances and RP sentences somewhat help

Results – Task 1A

Paper ID	MQ	Clair_UMich
C90_2039	0.235	0.635
C94_2154	0.288	0.536
E03_1020	0.239	0.478
H05_1115	0.350	0.375
H89_2014	0.332	0.546
J00_3003	0.196	0.559
J98_2005	0.101	0.344
N01_1011	0.221	0.498
P98_1081	0.200	0.367
X96_1048	0.248	0.535

Large deviation in scores, across topics, from both systems

Results – Task 2

Paper ID	MQ (using Task 1A MMR)
C90_2039	0.293
C94_2154	0.120
E03_1020	0.196
H05_1115	0.321
H89_2014	0.320
J00_3003	0.367
J98_2005	0.233
N01_1011	0.284
P98_1081	0.206
Average	0.260

ROUGE-L scores here measure overlap over the abstract since we did not have human summaries

Low scores could be due to deviation between summary of citances and the abstract of the paper

Errors – Task 1A

Citing text: “The line of our argument below follows a proof provided in... for the maximum likelihood estimator based on nite tree distributions.”

Clair_UMich

False negative: “We will show that in both cases the estimated probability is tight.”

MQ

Target text from RP: “*The work described here also makes use of hidden Markov model.*”

False positive: “The statistical methods can be described in terms of Markov models.”

Learning from the Pilot Task

- Offset mismatch between the text file and the XML that annotators used
 - Corpus sentence segmented and sentences assigned a sentence ID
- Problems with handling non-contiguous spans in Protégé
- Character offsets can be miscounted by different parsers
- Handling non-UTF8 characters

Limitations of this corpus

- No gold standard citation based summaries

- OCR errors: **Rolf Kümmerli,^{1,2} Andy**



Rolf K"ummerli,^{1,2}

- The use of “...” where text spans are snippets
- Errors in citation/reference offset numbers
- Different text encodings
- Errors in file construction
- **Small size of corpus!**

Acknowledgements

- NIST and Hoa Dang
- Lucy Vanderwende, MSR
- Horacio Saggion for detailed bug report on the corpus
- Rahul Jha (U. Mich, Ann Arbor), Prabha Yadav (U. Colorado, Boulder)
- All BiomedSumm track participants

Questions? Thank you!